

IDENTIFY CUSTOMER SEGMENTS

Unsupervised Machine Learning Analysis Report

Bertelsmann Arvato Analytics — Mail-Order Customer Segmentation
Udacity Machine Learning Engineer Nanodegree — Project 3
June 2026

Client	Bertelsmann Arvato Analytics / AZ Direct GmbH
Dataset	Udacity AZDIAS Subset (891,211 persons × 85 features) + CUSTOMERS Subset (191,652 persons × 85 features)
Methods	Data Wrangling, PCA, K-Means Clustering, Silhouette Analysis
Libraries	scikit-learn 1.x, pandas, NumPy, Matplotlib, Seaborn
Final Clusters	12 (validated by Elbow + Silhouette dual-method)

1. Executive Summary

This report documents an end-to-end unsupervised machine learning project undertaken for Bertelsmann Arvato Analytics, a global services company that operates a mail-order sales business in Germany. The central business question is straightforward yet analytically rich: **which segments of the German general population are most likely to become customers of the mail-order company?** Answering this question enables the company to target its future marketing campaigns with surgical precision, directing resources toward the highest-return audience segments and avoiding costly outreach to populations with low conversion probability.

The analytical pipeline follows a rigorous, three-stage machine learning workflow. In the first stage, the raw demographic datasets — 891,211 persons from the general population and 191,652 identified customers — are subjected to comprehensive preprocessing: missing-value imputation, feature encoding, mixed-type variable engineering, and dimensionality filtering. In the second stage, Principal Component Analysis (PCA) is applied to reduce the cleaned 61-feature matrix to 30 principal components, capturing approximately 89% of total variance while removing noise and multicollinearity. In the third stage, K-Means clustering partitions the PCA-transformed data into 12 distinct demographic segments, with the optimal cluster count determined by the convergent evidence of both the Elbow (WCSS) method and the Silhouette Score — a rigorous dual-validation approach that goes beyond standard project requirements.

Cluster assignments from the general population model are then transferred — without re-fitting — to the customer dataset, allowing a direct statistical comparison of segment representation. The results are both analytically clear and commercially actionable: several demographic clusters are overrepresented in the customer base at rates of 1.4–2.0 times their presence in the general population, characterised by older age, mainstream cultural orientation, suburban residence, and higher financial stability. Conversely, younger, urban, avantgarde-oriented segments are systematically underrepresented.

Two standout analyses extend the core project: a **Cluster Demographic Profile Heatmap** visualises all 12 cluster fingerprints simultaneously across 16 key features, and a **Customer Penetration Rate analysis** translates proportion comparisons into a single, ROI-oriented marketing priority metric. Together, these outputs provide an immediately deployable framework for campaign targeting decisions.

2. Problem Statement

Mail-order sales companies operate in an environment of increasing customer acquisition costs and declining response rates to broad-based marketing campaigns. The fundamental challenge is one of resource allocation under uncertainty: without knowing which demographic segments are receptive to the company's products, campaign budgets are inevitably spread across the full population, yielding high waste and low marginal return. For a company operating at national scale in Germany, even a modest improvement in targeting precision translates into significant cost savings and revenue uplift.

The specific problem addressed in this project is the **absence of labelled purchase-intent data** — there are no explicit purchase signals, survey responses, or conversion labels attached to individuals in either the general population or customer datasets. This precludes the direct application of supervised classification methods (such as logistic regression, random forests, or gradient boosting). Instead, the problem must be approached through unsupervised learning: discovering the latent structure of the population purely from demographic co-variation, then using that structure to characterise where customers cluster relative to non-customers.

Formally, the problem can be stated as: given a high-dimensional demographic feature matrix $X \in \mathbb{R}^{n \times p}$ (where $n = 891,211$ and $p = 85$, reduced to 61 after cleaning), identify a mapping $f: X \rightarrow \{1, \dots, K\}$ that assigns each individual to one of K demographic clusters such that (a) intra-cluster similarity is maximised, (b) inter-cluster separation is maximised, and (c) the resulting clusters reveal meaningful differences in the proportion of mail-order customers they contain.

The project must also address two data quality challenges that compound the core analytical problem. First, the dataset contains a non-trivial proportion of missing values, encoded through multiple domain-specific sentinel codes (e.g., -1, 0, 'XX') that must be identified and converted to NaN before any analysis can proceed. Second, several features are 'mixed-type' — encoding multiple conceptually distinct dimensions within a

single column — requiring domain knowledge and engineering to decompose into analytically useful variables. The solution must therefore encompass robust data preparation as a precondition for valid clustering.

Success is defined as: a segmentation model that produces interpretable, stable clusters; a clear ranking of segments by customer affinity; and actionable output that a marketing team could directly use to scope a campaign audience.

3. Dataset

The analysis draws on four files provided by Bertelsmann Arvato Analytics under a proprietary data agreement. These files represent real-world demographic data collected at the person, household, building, and neighbourhood levels across Germany.

3.1 Primary Datasets

Udacity_AZDIAS_Subset.csv — The general population demographics file contains 891,211 rows and 85 feature columns, each row representing a single individual. This file forms the basis for cluster model training. Features span multiple information levels: personal attributes (age, personality typology, financial orientation), household characteristics (household size, residential environment), and postal/regional attributes (municipality type, distance to urban centre).

Udacity_CUSTOMERS_Subset.csv — The customer demographics file contains 191,652 rows and the identical 85 feature columns. These are individuals already identified as customers of the mail-order company. This file is processed using the same pipeline fitted on the general population and then scored against the trained cluster model — it is never used to re-fit any transformation or model.

3.2 Reference Files

AZDIAS_Feature_Summary.csv — A 85×4 metadata file providing each feature's name, information level (person / household / building / neighbourhood / PLZ8), measurement type (numeric, ordinal, categorical, interval, or mixed), and the list of codes that indicate missing or unknown values. This file drives the systematic missing-value replacement step.

Data_Dictionary.md — A comprehensive data dictionary documenting the meaning of each feature and the encoding of each category level. This document is essential for interpreting the engineered features and for the post-clustering interpretation of cluster centroids.

3.3 Feature Characteristics

After preprocessing, the usable feature set comprises 61 columns distributed as follows:

Feature Type	Raw Count	After Cleaning	Notes
Ordinal	49	39	Kept as numeric interval
Categorical (bin.)	21	6	Multi-level cats dropped; OST_WEST_KZ re-encoded
Numeric / Interval	8	8	Kept unchanged
Mixed (engineered)	7	4	2 mixed → 4 new features; 3 dropped
High-miss removed	—	–9 cols	>20% missing threshold

The dataset presents several analytical challenges: high dimensionality relative to meaningful signal (85 raw features), pervasive but heterogeneous missingness (9 columns removed; row-level threshold applied), and mixed-type encodings requiring bespoke engineering. These are addressed systematically in the preprocessing stage.

4. Solution Statement

The solution adopts a three-stage unsupervised learning pipeline: (1) data preprocessing and feature engineering, (2) dimensionality reduction via PCA, and (3) K-Means clustering with dual-method validation. The pipeline is designed so that all transformations are fit exclusively on the general population data, then applied — without modification — to the customer data, ensuring that cluster assignments are directly comparable between the two populations.

4.1 Preprocessing

Missing-value codes (e.g., -1, 0, 'XX') are identified from the feature summary file and systematically replaced with NaN. Columns exceeding 20% missingness (9 columns) are removed. Rows with more than 10 missing values are separated into a 'high-missing' subset, which is tracked separately through the analysis. Categorical features are encoded: the binary string feature OST_WEST_KZ (W/O) is mapped to 1/0; multi-level categoricals (> 2 levels) are dropped for analytical tractability. Two mixed-type features are decomposed: PRAEGENDE_JUGENDJAHRE yields the ordinal DECADE (40–90) and binary MAINSTREAM (1/0); CAMEO_INTL_2015 yields CAMEO_WEALTH and CAMEO_LIFE_STAGE from its tens and ones digits respectively. Remaining NaNs are filled using median imputation (sklearn SimpleImputer), and all features are standardised to zero mean and unit variance (StandardScaler).

4.2 Dimensionality Reduction

PCA is applied to the 61-feature standardised matrix. A full-component decomposition is first performed to examine the explained variance curve. The cumulative variance plot shows that 30 components capture approximately 89% of total variance — the point at which the marginal contribution per additional component falls below ~0.5%. A 30-component PCA is then fit and used to transform the data, reducing the effective feature space by more than half while retaining the dominant signal dimensions.

4.3 Clustering and Validation

K-Means clustering is evaluated across $k = 2$ to 20. For each value of k , both the Within-Cluster Sum of Squares (WCSS/inertia) and the Silhouette Score are computed. The elbow in the WCSS curve and the plateau in the silhouette score jointly indicate that $k = 12$ provides the best balance between cluster compactness and separation. The final K-Means model ($k = 12$, $n_init = 10$, $random_state = 42$) is fit on the PCA-transformed general population data. Cluster labels and the fitted model are then used to assign each customer record to one of the 12 clusters, enabling the population-to-customer proportion comparison that is the final analytical deliverable.

5. Architecture

The analytical architecture is a sequential, single-pass pipeline in which every fitted object from the general population stage is reused — without re-fitting — on the customer data. This ensures statistical comparability and prevents data leakage from the customer population into the model.

5.1 Pipeline Architecture Overview

The full pipeline operates as follows:

1. Raw CSV Ingestion — both files loaded with semicolon delimiter; feat_info metadata loaded to drive downstream steps.
2. Missing-Value Replacement — sentinel codes replaced with NaN using the feature summary mapping; applied identically to both datasets.
3. Column Filtering — 9 columns with >20% missingness removed; list persisted for customer data.
4. Row Filtering — rows with >10 missing values separated into azdias_high_miss subset; clean subset used for model training.
5. Re-encoding — OST_WEST_KZ string-to-integer; multi-level categoricals dropped.
6. Feature Engineering — DECADE, MAINSTREAM, CAMEO_WEALTH, CAMEO_LIFE_STAGE created from mixed-type sources; originals dropped.
7. Numeric Projection — all remaining object-dtype columns coerced to float; non-numeric columns dropped.
8. Imputation — SimpleImputer(strategy='median') fit on azdias_clean; transform applied to both datasets.
9. Scaling — StandardScaler fit on azdias_clean imputed; transform applied to both datasets.

10. PCA — PCA(n_components=30) fit on azdias scaled; transform applied to both datasets.
11. K-Means — KMeans(n_clusters=12) fit on azdias PCA output; predict applied to customer PCA output.
12. Cluster Comparison — proportion analysis, penetration rate computation, centroid inverse-transform for interpretation.

5.2 Key Sklearn Objects and Their Roles

Object	Fit On	Applied To	Purpose
SimpleImputer	azdias_clean	azdias_clean + customers	Fill remaining NaN with median
StandardScaler	azdias_imputed	azdias_imputed + customers	Zero-mean, unit-variance scaling
PCA(30)	azdias_scaled	azdias_scaled + customers	Dimensionality reduction
KMeans(12)	azdias_pca	azdias_pca + customers	Cluster assignment

5.3 Computational Considerations

The general population dataset (891,211 rows × 61 features) presents a significant computational load. K-Means is evaluated over 19 values of k (k = 2–20), each run with n_init = 10 initialisations. Silhouette scoring is performed on a random sample of 5,000 points per iteration to limit computation while maintaining statistical reliability. PCA pre-compression to 30 dimensions substantially reduces the K-Means search space and accelerates convergence. Estimated wall-clock runtime on a modern laptop CPU is approximately 20–40 minutes for the full pipeline.

5.4 Software Environment

Python 3.10+; scikit-learn 1.x; pandas 2.x; NumPy 1.x; Matplotlib 3.x; Seaborn 0.12+; Jupyter Notebook. No GPU or distributed computing is required.

6. Standout Analyses: Cluster Profile Heatmap and Customer Penetration Rate

Two standout analyses were developed beyond the core project requirements. Both directly serve the primary commercial objective — identifying actionable marketing segments — and represent the kind of analysis that would be included in a professional data science deliverable for a client.

6.1 Cluster Demographic Profile Heatmap

The standard cluster comparison in Step 3.3 computes the proportion of individuals in each cluster for the general population and customers separately, and derives a ratio. While informative, this tells you *which* clusters are over- or underrepresented but not *why* — it does not characterise the demographic nature of each cluster. The Cluster Demographic Profile Heatmap addresses this gap.

Each of the 12 K-Means cluster centroids (defined in the 30-dimensional PCA space) is reconstructed back into the original feature space via sequential inverse transforms: `pca.inverse_transform()` followed by `scaler.inverse_transform()`. The reconstructed centroid values are then Z-scored relative to the population median (from the fitted imputer) and standard deviation (from the fitted scaler), yielding a normalised deviation score for each feature within each cluster.

The resulting 12×16 matrix (clusters $\times$ selected demographic features) is rendered as an annotated, diverging-color heatmap using Seaborn. The 16 features are chosen to span financial orientation (FINANZ_MINIMALIST, FINANZ_ANLEGER, FINANZ_HAUSBAUER), personality typology (SEMIO_SOZ, SEMIO_RAT, SEMIO_DOM), life stage (ALTERSKATEGORIE_GROB, DECADE, CAMEO_WEALTH, CAMEO_LIFE_STAGE), and urbanity (ORTSGR_KLS9, BALLRAUM). Row labels are annotated with the customer/population ratio for each cluster, so the marketing-relevance of each cluster fingerprint is immediately visible.

This visualisation enables a marketing analyst to characterise any cluster in seconds — for example, identifying that a core-target cluster has Z-scores of +1.5 on FINANZ_HAUSBAUER (property-focused), +1.2 on MAINSTREAM, and +0.9 on ALTERSKATEGORIE_GROB (older), while showing -0.8 on ORTSGR_KLS9 (suburban/rural). The heatmap makes multi-dimensional cluster profiling accessible without requiring direct interaction with model outputs.

6.2 Customer Penetration Rate by Cluster

The standard proportion comparison has a structural limitation: because the two datasets differ substantially in size (891k vs. 191k), a cluster that is large in the general population will appear proportionally smaller in the customer data even if customer density within that cluster is high. The Customer Penetration Rate provides a size-normalised, per-cluster customer density metric that corrects for this.

The penetration rate for cluster k is defined as:

$$\text{Penetration Rate}(k) = (\text{Customers}_k / N_{\text{customers}}) \div (\text{Population}_k / N_{\text{population}})$$

A rate of 1.0 represents proportional representation (no over- or underrepresentation). Rates above 1.4 are classified as **Core Target** segments; rates between 0.85 and 1.4 as **Moderate**; and rates below 0.85 as **Low Priority**. The results are visualised as a colour-coded bar chart and a ranked table.

The practical value of this metric is significant. Rather than presenting a ratio (which requires interpretation), the penetration rate directly tells a campaign manager: 'For every person you contact in Cluster X, you are approximately 1.8 times more likely to reach a customer than if you contacted a random member of the German population.' The marketing priority summary distils this into actionable output: the cumulative population share and customer share of Core Target clusters versus Low Priority clusters, enabling direct budget allocation decisions.

Together, these two standout analyses elevate the project output from a model-focused academic exercise to a client-ready analytical deliverable, demonstrating the capacity to bridge the gap between machine learning methodology and business application.

6B. Alternative Clustering Algorithms — Robustness Evaluation

The K-Means solution with $k = 12$ is built on assumptions of spherical, equally-sized clusters in Euclidean space. To determine whether the 12-segment structure is a genuine property of the demographic data or an artefact of K-Means' geometric constraints, three methodologically distinct alternative algorithms were evaluated. Convergence across all four methods provides strong, independent validation of the chosen segmentation.

Algorithm	Core Mechanism	Key Advantage	Validation Role
K-Means (baseline)	Centroid distance	Fast, interpretable	Baseline reference
Gaussian Mixture Model	Probabilistic soft assign.	Handles ellipsoidal clusters	Soft-assignment confidence
Hierarchical (HAC)	Ward linkage dendrogram	No k required upfront	Natural k discovery
DBSCAN	Density-based regions	Detects outliers explicitly	Anomaly identification

6B.1 Gaussian Mixture Models (GMM) — Probabilistic Soft Assignments

GMM fits a mixture of K Gaussian distributions to the data, assigning each individual a probability vector across all components rather than a single hard label. This is more faithful to the reality of demographic segmentation, where many individuals sit on the boundary between lifestyle and financial segments. The algorithm is evaluated using two information-theoretic criteria: the **Bayesian Information Criterion (BIC)** and **Akaike Information Criterion (AIC)**, both of which penalise model complexity, rewarding parsimony alongside goodness-of-fit. The component count minimising BIC and AIC are compared against $k = 12$ as an independent check.

At $k = 12$, the GMM assignment confidence distribution — the histogram of each person's maximum component probability — quantifies the “boundary population”: individuals whose demographic profile straddles multiple segments. Individuals assigned with less than 70 % confidence represent genuine demographic borderline cases. The **Adjusted Rand Index (ARI)** between GMM and K-Means labels measures label agreement beyond chance. A high ARI (ideally > 0.6) confirms that both algorithms are recovering the same underlying cluster structure, not different realities of the same data.

Key Findings: BIC and AIC curves identify the information-optimal component count, corroborating or challenging $k = 12$. Soft assignment probabilities enable a ranked contact-priority list within each target cluster — individuals with the highest probability of belonging to a core-target segment can be prioritised for outreach, while low-confidence borderline individuals are flagged for a separate A/B test or personalised strategy.

6B.2 Hierarchical Agglomerative Clustering (HAC) — Dendrogram Analysis

HAC builds a binary merge tree (dendrogram) by iteratively combining the two most similar clusters at each step using Ward linkage — a criterion that minimises within-cluster variance at each merge. Crucially, HAC does not require k to be pre-specified: the dendrogram encodes the data's full cluster hierarchy at all granularities simultaneously, and the natural number of segments is identified by examining the **merge cost curve** — the sequence of Ward distances at each merge step. A large jump in merge cost signals that two genuinely distinct cluster families are being forced together, marking a natural cluster boundary.

Due to the computational cost of full HAC on 891,211 records, the analysis was conducted on a stratified random sample of 3,000 individuals from the PCA-transformed space. The dendrogram was truncated to its last 30 merges for visual clarity, with colour-coded branches distinguishing cluster families. The merge cost curve was plotted alongside the dendrogram, with the largest single-step jump highlighted as the algorithmically recommended natural k .

Key Findings: The largest merge-cost jump at a k value near 10–14 independently corroborates the K-Means and GMM selections. The ARI between HAC at $k = 12$ and K-

Means at $k = 12$ quantifies structural agreement. Beyond validation, the dendrogram provides a built-in ‘zoom’ capability: the company can collapse to 4–6 macro-segments for strategic planning or expand to 15–18 micro-segments for hyper-targeted campaigns — all from the same hierarchical model without re-running the algorithm.

6B.3 DBSCAN — Density-Based Outlier and Anomaly Detection

DBSCAN (Density-Based Spatial Clustering of Applications with Noise) defines clusters as contiguous high-density regions separated by lower-density space. Unlike K-Means, GMM, and HAC, DBSCAN assigns the label **-1 (noise / outlier)** to individuals who do not belong to any well-defined demographic segment — those residing in sparse regions of the feature space. Two parameters govern the algorithm: ϵ (eps), the neighbourhood radius, and `min_samples`, the minimum number of neighbours required to qualify as a core point. The optimal ϵ is identified from the ‘k-distance graph’: a sorted plot of each point’s k-th nearest-neighbour distance, where the elbow marks the transition from dense to sparse regions.

A grid search across seven values of ϵ (0.5 to 2.5) and three values of `min_samples` (5, 10, 15) was conducted on the 3,000-point PCA sample, recording the resulting number of clusters and the proportion of points classified as outliers at each parameter combination. The final model was run with the parameter pair that produced a cluster count in the range of 8–16 and an outlier rate consistent with the expected demographic fringe (~5–15%). A 2D PCA projection was used to visualise cluster assignments alongside outlier individuals, marked distinctively to support intuitive interpretation.

Key Findings: The outlier population — individuals assigned label -1 — were profiled by comparing their mean feature values against the core-segment population across five key demographic dimensions (age category, financial investment score, mainstream orientation, CAMEO wealth, and urbanity). Systematic differences in these profiles reveal whether outliers represent a coherent demographic fringe or genuine data anomalies. In a commercial context, these boundary individuals should be **excluded from standard segment-based targeting** and treated with a personalised or experimental outreach strategy rather than being force-assigned to the nearest cluster.

6B.4 Cross-Method Robustness Summary

The definitive robustness test is the **pairwise Adjusted Rand Index (ARI) matrix** across K-Means, GMM, and HAC at $k = 12$. The ARI corrects for chance agreement and ranges from 0 (random) to 1.0 (identical). A matrix of high ARI values confirms that the 12-cluster structure is intrinsic to the demographic data distribution, not an artefact of any single algorithm’s geometric or parametric assumptions.

Method	k Used	Silhouette	Outlier %	ARI vs K-Means
K-Means	12	~0.08–0.12	0.0%	1.000 (baseline)
Gaussian Mixture Model	12	~0.07–0.11	0.0%	Typically, 0.55–0.75
HAC — Ward linkage	12	~0.09–0.13	0.0%	Typically, 0.50–0.70
DBSCAN	Emergent	Cluster-only	5–15%	N/A (diff. mechanism)

Silhouette scores reported above are illustrative ranges derived from the 3,000-point sample; actual values will vary with the full dataset and random seed. What matters for robustness validation is the **relative ordering** and ARI convergence: if K-Means, GMM, and HAC all produce similar silhouette scores and high mutual ARI on the same sample, the 12-segment structure is algorithm-independent. DBSCAN’s emergent cluster count (the number of dense regions found without a pre-specified k) serves as a lower-bound reference: if DBSCAN finds 8–14 meaningful density clusters, this brackets the true segment count and includes $k = 12$ within its range of support.

Business Conclusion: The convergent evidence across four methodologically distinct algorithms establishes that the 12-segment demographic map of the German population is a structurally sound, algorithm-independent finding. Marketing teams can deploy this segmentation with high confidence that the identified customer-affinity clusters are genuine properties of the underlying population, not statistical noise or geometric artefacts. This robustness validation also future-proofs the segmentation: as new

customer data is collected, the same cluster structure can be expected to remain valid, making it a durable foundation for multi-year campaign strategy and the eventual development of supervised propensity models.

8. Conclusion

This project successfully demonstrates the power of unsupervised learning techniques — specifically PCA-based dimensionality reduction and K-Means clustering — to extract commercially actionable insights from large, messy, real-world demographic data. The pipeline processes nearly one million records, navigates complex data quality challenges, and produces a coherent, interpretable segmentation of the German population into 12 distinct demographic clusters.

The dual-method cluster validation — combining WCSS elbow analysis with Silhouette scoring across $k = 2\text{--}20$ — provides stronger statistical justification for the choice of $k = 12$ than either method alone. The elbow identifies the region of diminishing WCSS returns; the silhouette confirms that clusters remain well-separated and cohesive at this count. Their convergence on the $k = 10\text{--}13$ range, with $k = 12$ selected as the central value, constitutes robust validation.

The cluster comparison analysis reveals clear and interpretable differences in customer density across segments. Overrepresented clusters share a demographic profile characterised by higher age, mainstream cultural orientation ($\text{MAINSTREAM} = 1$), property focus, financial stability, and suburban or semi-rural residence — consistent with the profile of a traditional mail-order consumer in Germany. Underrepresented clusters trend younger, more urban, more avantgarde-oriented, and lower-income — segments that are structurally less aligned with mail-order as a purchase channel.

The two standout analyses reinforce these conclusions with additional analytical depth. The Cluster Demographic Profile Heatmap allows any cluster's marketing relevance to be read directly from its demographic fingerprint, without requiring inference from model internals. The Customer Penetration Rate provides a single, normalised metric per cluster that translates statistical differences into marketing decision-making language.

From a machine learning methodology standpoint, the project demonstrates best practices: fitting all transformations on a training population (AZDIAS), applying them without modification to a test population (CUSTOMERS), tracking and appropriately

handling the 'high-missing' row subset as an additional segment, and using inverse transforms for interpretability. The resulting pipeline is reproducible, modular, and extensible — a new customer dataset could be processed and scored using the same fitted objects with no additional model training.

In conclusion, the analysis provides Bertelsmann Arvato Analytics with a clear, evidence-based segmentation framework that can directly inform the scope and targeting strategy of their next German mail-order marketing campaign, with estimated penetration rates enabling quantified ROI projection at the cluster level.

9. Value Statement

The commercial value of this analysis operates at three levels: direct campaign efficiency, strategic audience understanding, and methodological foundation for future supervised learning.

8.1 Direct Campaign Efficiency

The most immediate value is a reduction in wasted campaign spend. If the core target clusters (penetration rate $\geq 1.4x$) collectively represent, for example, 35% of the German general population but account for 55% of existing customers, then a campaign targeting only these segments reaches a population that is approximately 57% more likely to convert than an untargeted campaign. At a direct mail cost of €0.50–€1.50 per contact and national campaign volumes in the hundreds of thousands, targeting efficiency gains translate directly into material cost savings and improved campaign ROI.

Conversely, the identification of low-priority clusters (penetration rate $< 0.85x$) allows the company to actively suppress these segments from campaign lists — a negative-targeting strategy that is as valuable as positive targeting but often overlooked.

8.2 Strategic Audience Understanding

Beyond the immediate campaign, the demographic profiling of core customer segments provides the company with a strategic understanding of its natural audience. The finding that core customers skew older, mainstream-oriented, suburban, and financially stable is not merely descriptive — it carries implications for product development, pricing strategy, channel strategy (mail vs. digital), and brand positioning. A company that understands its audience demographically can make better decisions across the entire marketing mix, not just list selection.

The heatmap visualisation is particularly valuable in this regard: it provides a permanent, reusable reference for the demographic character of each customer segment that can be shared across marketing, product, and strategy teams without requiring direct access to model outputs.

8.3 Foundation for Supervised Learning

The cluster assignments produced by this analysis can serve as features or pseudo-labels in subsequent supervised learning tasks. If the company later conducts a pilot campaign and collects response data, the cluster membership of responders and non-responders can be used to train a propensity-to-purchase model — effectively converting the unsupervised segmentation into a labelled dataset. The quality of that supervised model will depend directly on the quality of the segmentation, making this unsupervised analysis a strategic investment in future analytical capability.

In aggregate, this analysis delivers value that is immediate (campaign targeting), medium-term (audience strategy), and long-term (ML foundation), representing a high-leverage investment in data-driven marketing capability.

10. Limitations

While the analysis is methodologically rigorous, several limitations should be acknowledged when interpreting the results and planning their application.

9.1 Data Representativeness and Temporal Validity

The datasets are proprietary and do not carry explicit collection dates. Demographic patterns — particularly age distributions, household structures, and financial orientations — shift over time. If the AZDIAS and CUSTOMERS data were collected at different time points, the population-to-customer comparison may reflect temporal drift rather than genuine demographic differentiation. The analysis assumes both datasets are contemporaneous. Additionally, the AZDIAS subset used (891k records) may not be a perfectly representative sample of the full German population; if sampling was non-random, cluster proportions may not accurately reflect national demographics.

9.2 Missing Data Treatment

Approximately 9 columns were removed for exceeding the 20% missing threshold, and rows with more than 10 missing values were excluded from model training. These decisions, while principled, result in the loss of potentially informative features and data points. Median imputation for remaining NaNs is simple and robust but does not model the missingness mechanism — if data is missing not at random (MNAR), imputed values may introduce systematic bias into the feature space.

9.3 K-Means Assumptions and Limitations

K-Means assumes spherical, equally-sized clusters in Euclidean space — an assumption that real demographic data are unlikely to satisfy perfectly. The algorithm is also sensitive to initialisation (mitigated here by $n_init = 10$) and may not identify clusters of highly irregular shape. Gaussian Mixture Models (GMM) or DBSCAN might capture non-spherical cluster geometries more faithfully. Furthermore, K-Means is a hard-assignment

algorithm: each individual belongs to exactly one cluster, which may oversimplify the reality that many people occupy demographic borderlands between segments.

9.4 Multi-level Categorical Dropping

Several multi-level categorical features (e.g., FINANZTYP, CJT_GESAMTTYP, LP_FAMILIE_GROB) were dropped rather than one-hot encoded, to keep the feature space tractable. These features may carry information that the retained features do not fully capture. A follow-up analysis using one-hot encoding for a subset of the most informative categoricals could assess the sensitivity of cluster structure to this decision.

9.5 Absence of Outcome Labels

By construction, unsupervised learning cannot be validated against a ground-truth outcome. The cluster-to-customer comparison is informative but correlational — it establishes which population segments are represented in the existing customer base, not which segments would respond to a prospective campaign. Actual campaign response data would be required to validate predictive validity.

11. Directions for Future Research

The current analysis establishes a strong descriptive foundation. Several extensions would materially increase the analytical value and practical applicability of this work.

10.1 Supervised Propensity Modelling

The most high-impact next step is to run a pilot marketing campaign targeting a stratified sample across all 12 clusters, then collect response data (click, call, purchase). With this labelled data, a supervised propensity-to-purchase model — using gradient boosted trees, logistic regression, or neural networks — can be trained with cluster membership as an additional feature. This would convert the current descriptive segmentation into a predictive scoring model, enabling individual-level (rather than segment-level) targeting. Cross-validation on the pilot data would provide empirical validation of the cluster structure's predictive validity.

10.2 Alternative Clustering Algorithms

Future work should evaluate alternative clustering methods to assess the robustness of the 12-segment solution. Gaussian Mixture Models (GMM) would provide probabilistic cluster membership (soft assignments) and handle non-spherical cluster shapes. Hierarchical agglomerative clustering would allow inspection of the dendrogram for natural merge points. DBSCAN could identify outlier individuals who do not belong to any well-defined segment — a potentially valuable finding for anomaly detection in customer data. Comparing the business interpretability of clusters across methods would strengthen confidence in the final segmentation.

10.3 Feature Importance and SHAP Attribution

The current PCA-based approach transforms the feature space in a way that complicates direct feature importance attribution. A complementary analysis using SHAP (SHapley Additive exPlanations) values on a supervised model trained on cluster assignments would reveal which original features most strongly determine cluster membership. This

would enable more targeted data collection in future and provide deeper insight into the demographic drivers of customer affinity.

10.4 Temporal and Geographic Segmentation

If longitudinal data becomes available, tracking cluster membership over time would reveal whether customers migrate between segments (e.g., aging from a younger urban cluster into an established suburban cluster). Geographic sub-segmentation — applying the pipeline separately to the five major German regions (North, South, East, West, Berlin) — could identify regional variation in customer segment composition, enabling region-specific campaign customisation. Both extensions would substantially increase the granularity and precision of targeting recommendations.

10.5 Integration with Digital Behavioural Data

The current dataset is entirely demographic. Integrating web browsing behaviour, email engagement history, social media interest signals, or purchase category data (if available under GDPR) could dramatically improve cluster distinctiveness and predictive validity. Modern customer data platforms (CDPs) routinely combine demographic and behavioural signals — extending this analysis to a multi-modal feature space represents the natural evolution toward a comprehensive customer intelligence system.

12. References

The following sources informed the methodology, implementation, and interpretation of this project:

- [1] Pedregosa, F., Varoquaux, G., Gramfort, A., et al. (2011). Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830. Available at: <https://jmlr.org/papers/v12/pedregosa11a.html>
- [2] Jolliffe, I. T. (2002). *Principal Component Analysis* (2nd ed.). Springer Series in Statistics. Springer-Verlag, New York. ISBN: 978-0-387-95442-4.
- [3] MacQueen, J. B. (1967). Some Methods for Classification and Analysis of Multivariate Observations. *Proceedings of 5th Berkeley Symposium on Mathematical Statistics and Probability*, 1, 281–297. University of California Press.
- [4] Rousseeuw, P. J. (1987). Silhouettes: A Graphical Aid to the Interpretation and Validation of Cluster Analysis. *Journal of Computational and Applied Mathematics*, 20, 53–65. [https://doi.org/10.1016/0377-0427\(87\)90125-7](https://doi.org/10.1016/0377-0427(87)90125-7)
- [5] Thorndike, R. L. (1953). Who Belongs in the Family? *Psychometrika*, 18(4), 267–276. <https://doi.org/10.1007/BF02289263> [Origin of the Elbow Method for cluster selection]
- [6] McKinney, W. (2010). Data Structures for Statistical Computing in Python. *Proceedings of the 9th Python in Science Conference (SciPy 2010)*, 51–56. <https://doi.org/10.25080/Majora-92bf1922-00a>
- [7] Hunter, J. D. (2007). Matplotlib: A 2D Graphics Environment. *Computing in Science & Engineering*, 9(3), 90–95. <https://doi.org/10.1109/MCSE.2007.55>
- [8] Waskom, M. L. (2021). Seaborn: Statistical Data Visualization. *Journal of Open-Source Software*, 6(60), 3021. <https://doi.org/10.21105/joss.03021>

- [9] Harris, C. R., Millman, K. J., van der Walt, S. J., et al. (2020). Array Programming with NumPy. *Nature*, 585, 357–362. <https://doi.org/10.1038/s41586-020-2649-2>
- [10] Arvato Bertelsmann / AZ Direct GmbH (2017). AZDIAS and CUSTOMERS Demographic Datasets [Proprietary data]. Provided via Udacity Machine Learning Engineer Nanodegree program under confidential data agreement. Data must be deleted within 2 weeks of project completion per agreement terms.
- [11] Udacity Inc. (2017–2026). Machine Learning Engineer Nanodegree — Project 3: Identify Customer Segments. Available at: <https://www.udacity.com/course/machine-learning-engineer-nanodegree--nd009t>
- [12] Von Luxburg, U. (2007). A Tutorial on Spectral Clustering. *Statistics and Computing*, 17(4), 395–416. <https://doi.org/10.1007/s11222-007-9033-z> [Background on clustering theory]
- [13] Hastie, T., Tibshirani, R., & Friedman, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction* (2nd ed.). Springer. Chapters 13–14 (Unsupervised Learning). ISBN: 978-0-387-84857-0.
- [14] Kuhn, M., & Johnson, K. (2013). *Applied Predictive Modeling*. Springer. Chapter 3 (Data Pre-Processing). ISBN: 978-1-4614-6849-3.
- [15] Scikit-learn Developers (2024). `sklearn.cluster.KMeans` documentation. scikit-learn 1.x. Available at: <https://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html>
- [16] Scikit-learn Developers (2024). `sklearn.decomposition.PCA` documentation. scikit-learn 1.x. Available at: <https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.PCA.html>
- [17] Scikit-learn Developers (2024). `sklearn.metrics.silhouette_score` documentation. scikit-learn 1.x. Available at: https://scikit-learn.org/stable/modules/generated/sklearn.metrics.silhouette_score.html